


Best measured results near half the dense computation

VideoMAE-B backbone

49.3% fewer FLOPs

96.4% of dense mAP retained




Cross: 68.58% mAP | 4,094 GFLOPs

Absolute mAP difference: -2.55 pp; selected EMA epoch 10

VideoMAE-S backbone

47.8% fewer FLOPs

93.6% of dense mAP retained



Cross: 64.57% mAP | 1,226 GFLOPs

Absolute mAP difference: -4.44 pp; selected EMA epoch 20



Zero-based bars; grey bars are the dense reference (100%). All absolute mAP values and gaps remain visible.
Best observed checkpoints below 55% dense cost; THUMOS14, full 768-candidate window, 2 MAC. Comparison is within backbone.

Best measured checkpoints at the same K384 observation budget

Cross versus the latest corrected BMCR peaks: B +0.44 pp; S +0.74 pp

VideoMAE-B / K=384 observations


VideoMAE-S / K=384 observations




Same observation budget, not identical total FLOPs or training recipes. Corrected BMCR costs are not imputed. Cross includes decoder cost; its audited FLOPs appear on the main page. Historical S accuracy is retained without inventing a cost point.

Compute-performance distribution of H65/BMCR routes

(a) All measured operating points



(b) Lower-compute region




- Cross
 Spatial
 Interpolation / controls
 Original-head control
 Observed envelope*
 Best measured Cross EMA
- Joint T/D/S
 TCN
 Depth
 Feature loss


47 complete evaluations + 4 fixed references; THUMOS14, 211 videos / 792 windows, mAP@0.3:0.1:0.7.
 FLOPs: full-window operator profiles (2 MAC), not video totals. *Envelope mixes models/checkpoints; no fitted points.

Measured budget curves from the same checkpoint

(a) VideoMAE-S - J01, epoch 5 EMA




(b) VideoMAE-B - J01, epoch 5 EMA




Each backbone uses one checkpoint and all 8 measured T/D/S settings; each curve changes T only (384 -> 768).
Segments connect observations; no fitted or extrapolated points. Green brackets compare measured FLOPs, not runtime speed.

Performance spread across checkpoints at fixed computation

(a) VideoMAE-S



(b) VideoMAE-B



Fixed-compute checkpoint trajectories, not budget curves or confidence intervals; GFLOPs are shown in the legends.
Feature-loss-only retains GT boundary weighting. Epochs denote additional decoder training on a fixed anchor.